

Dropout: A Simple Way to Prevent Neural Networks from Overfitting

Srivastava et al

10.10.17

Summary by Justin Chen

Summary

Introduction:

- Dropout technique addresses better generalization for large, over-parameterized neural networks
- Technique **randomly** drops neurons during **training**
- Trains an ensemble of “**thinned**” networks and has the approximate effect of averaging predictions of ensemble during inference
- They focus on regularizing **fixed-sized models** by approximating an equally weighted geometric mean of the predictions from an ensemble of exponentially mean trained models that all share the same set of parameters
- **Naive approach** to averaging exponentially many models would require finding unique set of hyperparameters, unique dataset, training each individual network, and inferencing all networks simultaneously which is infeasible.
 - Dropout circumvents the aforementioned issues while also improving generalization
- All connections to dropped out neuron are also dropped during training
- Neurons are dropped with probability P
 - In practice, neurons in the input layer are dropped with much smaller probability, less than 0.8 and for hidden units and output units P can be fixed to 0.5
 - P can also be adjusted using validation set
- Dropout essentially samples subnetworks during training
- Original network with n neurons can also be thought of as having 2^n subnetworks
- All subnetworks share the same parameters
 - Total number of parameters is $O(n^2)$
 - Remember each row in a weight matrix is a neuron
- Training with dropout effectively
 - Trains at most 2^n networks, some aren't sampled during training

- During inference, weights of dropped out neurons are scaled down by probability P of it being dropped out during training
 - Neurons are **never** dropped out during inference. All neurons that were dropped out during training are **kept during inference** along with all other neurons!
 - Note that in this summary, P is the probability of being dropped and the paper refer to P as the probability of being retained
 - I think referring to the dropout probability in this way is more natural
 - P here is $1-P$ in the paper.
 - Scaling by P ensures that the expected output of that neuron is the same during training and inference
- They found that dropout significantly **decreases** generalization **error** compared to previous regularization techniques
 - Note: this paper was published in 2014
- Dropout encourages sparse activation during inference
- Dropout can be thought of as regularizing by adding noise to neurons
- “minimize loss function stochastically under noise distribution.”
 - Can be interpreted as minimizing expected loss function
- Does not require dimensionality reduction to perform well
- Gaussian noise can perform better than Bernoulli
- Drawbacks:
 - 2-3 times slower training time compared to same architectures without dropout
 - Noisy parameter updates

Motivation:

- Sexual reproduction in evolution
 - Half of a parent's genes are dropped (*innuendos - giggity*.)
 - Asexual reproduction should work better since the genes have evolved (optimized) to work together
 - Sexual reproduction breaks out those co-adaptations, which you'd assume would decrease fitness, yet sexual reproduction has evolved the most advanced organisms
 - “Mix-ability” - ability to be useful alone or work with other random genes
 - This may be the desired property of natural selection not individual fitness (hypothesis)
 - Allows new genes to spread throughout population and reduce complexity that would ultimately reduce fitness



- Breaking co-adaptation “makes each hidden unit more robust and drive it towards creating useful features on its own without relying on other hidden units to correct its mistakes.”

Model Description:

$$\begin{aligned}
 r_j^{(l)} &\sim \text{Bernoulli}(p), \\
 \tilde{y}^{(l)} &= r^{(l)} * y^{(l)} \\
 z_i^{(l+1)} &= w_i^{(l+1)} \tilde{y}^{(l)} + b_i^{(l+1)} \\
 y_i^{(l+1)} &= f(z_i^{(l+1)})
 \end{aligned}$$

- In the equations above, everything is a vector and $*$ means element-wise multiplication
- Essentially each weight has an extra Bernoulli coefficient (either 0 or 1)
- Derivative of loss function only backpropagated through subnetwork (neurons that were not dropped)

Backpropagation:

- Train as usual, but sample a subnetwork for each mini-batch
- “Gradients for **each parameter** are **averaged** over the training cases in **each mini-batch**”
- Dropped parameters receive **zero gradients**
- Combining dropout with max-norm regularization, decaying learning rate, and high momentum further improve generalization