


The Lottery Ticket Hypothesis: Training Pruned Neural Networks

Jonathan Frankle, Michael Carbin

**MACHINE
INTELLIGENCE
COMMUNITY**

Devin de Hueck
Mar. 29, 2018

Pruning Neural Networks



The Hypothesis

“Training succeeds for a given network if one of its subnetworks (a “winning ticket”) has been randomly initialized such that it can be trained in isolation to high accuracy in at most the number of iterations necessary to train the original network.”



Research Questions

“How effectively do winning tickets train in comparison to the original network and to randomly sampled networks of similar size?”

“How big are winning tickets relative to the size of the original network?”

“How sensitive are our results to particular pruning strategies?”



Extracting a “Winning Ticket”

1. Randomly Initialize a neural network
2. Train the network until convergence
3. Prune
4. Reset weights of the remaining portion to original initialization



Learning the XOR function - Results



10 Units		8 Units		6 Units		4 Units		2 Units	
DB	ZL	DB	ZL	DB	ZL	DB	ZL	DB	ZL
98.5	92.9	96.8	87.5	92.5	76.8	78.3	55.3	49.1	17.6

DB: Decision Boundary
ZL: Zero loss

Pruning Strategy	10 Units		4 Units (Pruned)		2 Units (Pruned)	
	DB	ZL	DB	ZL	DB	ZL
One-shot Product	99.2	93.3	98.0	90.3	82.4	65.3
Input Magnitude	98.9	93.5	97.9	92.2	83.8	76.5
Output Magnitude	99.0	93.3	96.9	85.9	78.6	56.1
Product	98.5	92.9	97.6	90.3	91.5	79.4




Iterative Pruning

MNIST - One Shot Pruning



Each curve represents the average performance of five trials in which networks were pruned to the stated size

MNIST - One Shot Pruning - Control Experiment 1

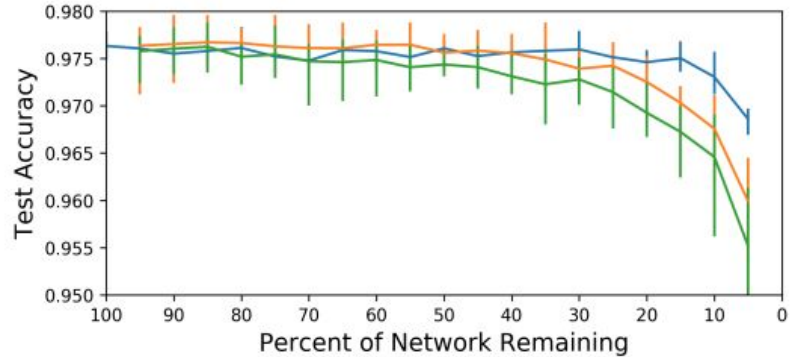
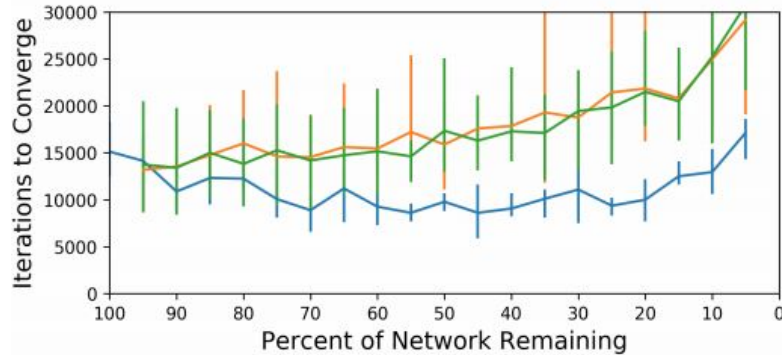


Blue curve: average of 5 winning tickets

Multi-colored curves: the controls for each of the 5 trials



MNIST - One Shot Pruning - Control Experiment 2



Blue Curve: average of 5 winning tickets

Orange Curve: results from control experiment 1



Green Curve: results from control experiment 2

MNIST - Iterative Pruning

1. Randomly Initialize the network
2. Train the network until convergence
3. Prune 20% from each hidden layer
4. Reset weights to original initialization
5. Repeat steps 2 through 4 until network is pruned to desired size





MNIST - Iterative Pruning - Results



Orange Curve: One-shot pruning results

Blue Curve: Iterative pruning results

MNIST - Iterative Pruning - Control Experiments



Blue Curve: average of five winning tickets pruned iteratively

Red Curve: one-shot pruning

Orange Curve: results of control experiment 1


Green Curve: results of control experiment 2




Examining Winning Tickets - Initializations

% of network
remaining:


100%




51.2%



16.8%

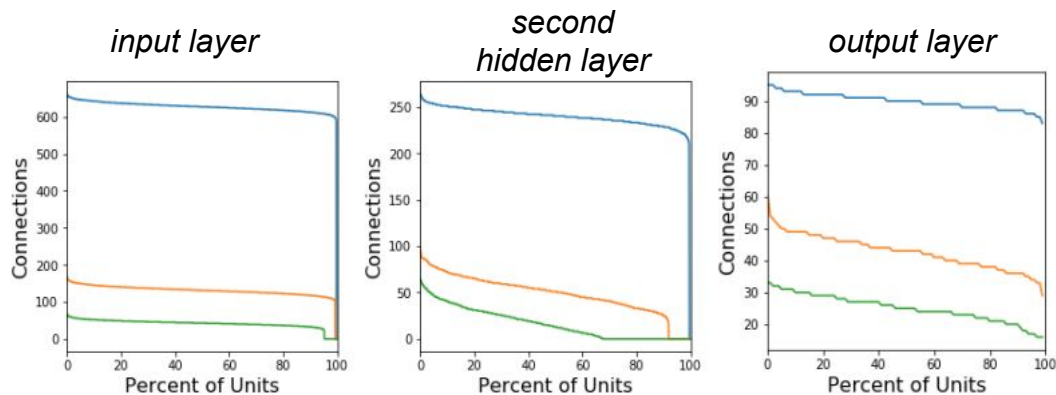


5.5%



Blue curve: first hidden layer
Orange curve: second hidden layer
Green curve: output layer

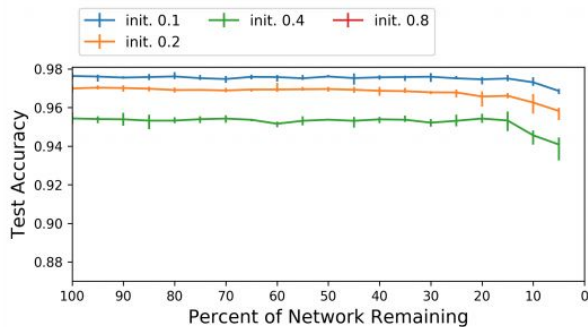

Examining Winning Tickets - Architecture





Blue Curve: iteratively pruned to 80%
Orange Curve: iteratively pruned to: 16.8%
Green Curve: iteratively pruned to 5.5%

Exploring MNIST Parameters - Initialization

STD ≥ 0.1

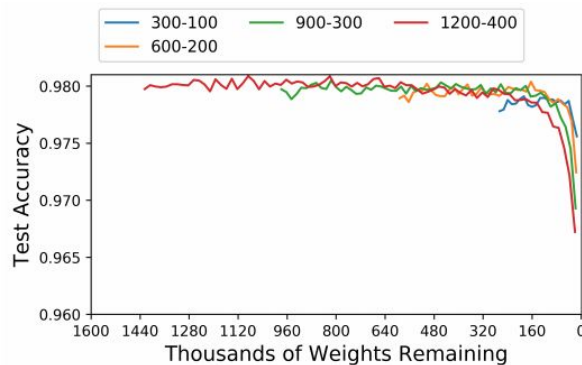



STD ≤ 0.1

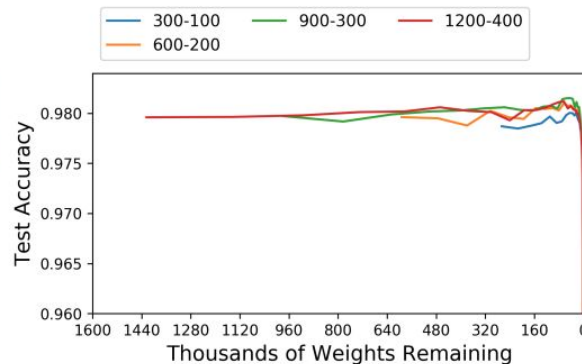



Exploring MNIST Parameters - Network Size


One-shot Pruning:



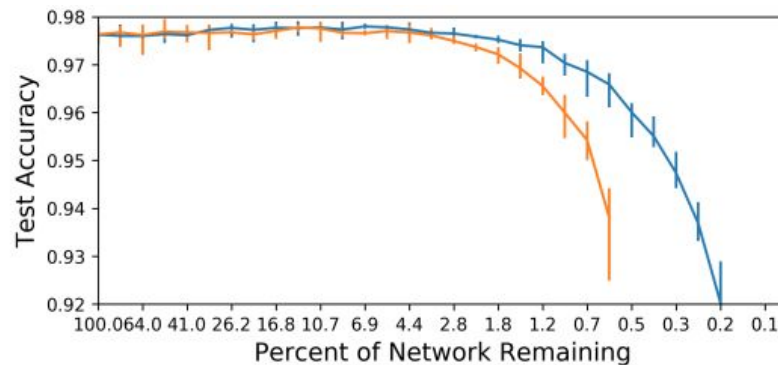

Iterative Pruning:



Exploring MNIST Parameters - Iterative Pruning Rates



Exploring MNIST Parameters - Weight Resetting



Blue: This papers strategy

Orange: Previous research strategy



Limitations

- Small Examples
- No Theoretical Analysis
- No Useful Strategies Devised

Conclusion and What's Next

- XOR
- MNIST
- Future Research Directions
 - Larger Examples
 - Understanding Winning Tickets
 - New Strategies



References and Further Reading

1. Han, Song, et al. **"Learning both weights and connections for efficient neural networks."** Advances in neural information processing systems. 2015.
2. Srivastava, Nitish, et al. **"Dropout: A simple way to prevent neural networks from overfitting."** The Journal of Machine Learning Research 15.1. 2014.
3. Han, Song, et al. **"Deep compression: Compressing deep neural networks with pruning, trained quantization and huffman coding."** arXiv preprint arXiv:1510.00149. 2015.
4. Hinton, Geoffrey E., et al. **"Improving neural networks by preventing co-adaptation of feature detectors."** arXiv preprint arXiv:1207.0580. 2012.
5. Lin, Henry W., et al. **"Why does deep and cheap learning work so well?"** arXiv preprint arXiv:1608.08225. 2017.

